

Willow M. Ahrens

ASSISTANT PROFESSOR AT GEORGIA TECH · COMPILERS, DATA STRUCTURES, ALGORITHMS

Office 3144, Klaus Advanced Computing Building, Georgia Tech

□ (+1) 505-412-5239 | □ ahrens@gatech.edu | □ willowahrens.io | □ willow-ahrens | □ 0QFbKQ4AAAAJ | □ 0000-0002-4963-0869

Education

Georgia Institute of Technology

Atlanta, GA

ASSISTANT PROFESSOR

Aug. 2025 - Present

Massachusetts Institute of Technology

Cambridge, MA

POSTDOCTORAL RESEARCHER. COMPUTER SCIENCE, PI: SAMAN AMARASINGHE

Sep. 2024 - Jul 2025

- Productionized Ph.D. research for deployment in real-world python ecosystems, integrating research to eventually include as a backend for pydata/sparse, in the default conda installation.
- Research conducted under a DARPA SBIR grant.
- Published 5 papers in top-tier conferences and journals, including OOPSLA, CGO, SIGMOD, and ASPLOS.
- Developed Finch.jl programming language and compiler for sparse and structured arrays.

Massachusetts Institute of Technology

Cambridge, MA

PH.D. COMPUTER SCIENCE, GPA: 4.9 / 5.0, ADVISOR: SAMAN AMARASINGHE

Sep. 2016 - Sep 2024

- Collaborated on intermediate languages, cost models, compiler passes, and algorithms for state-of-the-art research projects.
- Published 6 papers in top-tier conferences and journals, including PLDI, CGO, TOMS, and IPDPS.
- Presented at 15+ conferences, workshops, and research groups in academia and industry.
- Advised 4 undergraduates and 3 masters students. Proposed projects and provided weekly feedback. One student published in SPAA and another in CGO.
- Developed Finch.jl programming language and compiler for sparse and structured arrays.
- Discovered compiler algorithms to automatically adapt programs to input properties.

University of California, Berkeley

Berkeley, CA

BS IN EECS, MINOR IN MATH, GPA: 3.8 / 4.0, ADVISOR: JAMES DEMMEL

Sep. 2012 - May 2016

- Berkeley Benchmarking and Optimization Group
- Developed a reproducible linear algebra library in C, ReproBLAS.

Selected Publications

Finch: Sparse and Structured Tensor Programming with Control Flow

OOPSLA 2025

Apr. 2025

W. AHRENS, T. F. COLLIN, R. PATEL, K. DEEDS, C. HONG, AND S. AMARASINGHE.

- Built the Finch array programming language and compiler.
- Finch is the first compiler to automatically specialize flexible control flow to diverse array data structures.
- Supports a familiar programming language of loops, statements, ifs, breaks, etc., over a wide variety of array structures, such as sparsity, run-length-encoding, symmetry, triangles, padding, or blocks. Finch reliably utilizes the key properties of structure, such as structural zeros, repeated values, or clustered non-zeros.
- The first compiler to support convolution over sparse arrays, as well as worst-case optimal joins and variable-width blocks.
- Demonstrated dramatic speedups in operations such as SpMV and SpGEMM, image processing, and graph analytics.

Looplets: A Language for Structured Coiteration.

CGO 2023

Feb. 2023

W. AHRENS, D. DONENFELD, F. KJOLSTAD, AND S. AMARASINGHE.

- Published in Proceedings of the 21st ACM/IEEE International Symposium on Code Generation and Optimization, in CGO 2023.
- Proposed the Looplets language, the key IR behind the Finch tensor compiler. Allows flexible iteration over a wide variety of sparse and structured arrays.

Autoscheduling For Sparse Tensor Algebra With An Asymptotic Cost Model.

PLDI 2022

Jun. 2022

W. AHRENS, F. KJOLSTAD, AND S. AMARASINGHE.

- Published in Proceedings of the 43rd ACM SIGPLAN International Conference on Programming Language Design and Implementation.
- Discovered an asymptotic notation for the runtime of sparse tensor programs.
- Built the first asymptotic autoscheduler for sparse tensor compilers.

Tensor Algebra Compilation with Workspaces.

CGO 2019

2019

F. KJOLSTAD, W. AHRENS, S. KAMIL, AND S. AMARASINGHE.

- Published in 2019 IEEE/ACM International Symposium on Code Generation and Optimization (CGO), 2019, pp. 180-192.

A Fill Estimation Algorithm for Sparse Matrices and Tensors in Blocked Formats.

IPDPS 2018

2018

W. AHRENS, H. XU, AND N. SCHIEFER.

- Published in 2018 IEEE International Parallel and Distributed Processing Symposium (IPDPS), 2018, pp. 546-556.

Teaching

Georgia Tech Course CS8803-DSL (Domain Specific Languages for HPC)

Atlanta, GA

INSTRUCTOR

- Designed CS8803-DSL, a graduate-level seminar course on domain-specific languages for high-performance computing.
- Taught the basics of compiler construction and language design as it relates to the challenges of HPC and performance engineering.
- Covered topics such as architectural modelling, optimization techniques, and program analysis, through the context of contemporary DSLs and research in the field.

MIT Course 6.1200 (Mathematics For Computer Science)

Boston, MA

TEACHING ASSISTANT

- Taught 6.1200 (formally 6.042), a proof-based course designed to teach the fundamentals of algorithmic thinking in computer science, with attention given to concepts such as induction, asymptotic analysis, graphs, and probability.
- Led two discussion sections with 30 students each, covering example problems and their solutions.
- With 2 other TAs, staffed the last in-person office hours before homework was due each week, with attendance regularly exceeding 40 students requesting individual attention.

MIT Glass Lab

Cambridge, MA

GLASSBLOWING INSTRUCTOR

Feb 2019 – Present

- Supervised pairs of beginner students one at a time for weekly two-hour sessions.
- Ensured student safety in their first experiences with handling 2400 °F glass in a crowded hot shop.
- Explained critical techniques in glassblowing, including gathering, marvering, blocking, and blowing.

Center for Access to Engineering Excellence

Berkeley, CA

TUTOR

Jan 2014 – May 2014

- Tutored groups of around 2-5 students at a time in lower-division Computer Science, Math, and Physics courses.

Honors

2017-2021 **CSGF Fellow**, DOE Computational Science Graduate Fellow

Washington, D.C.

2016 **Warren Y. Dere Design Award**, UC Berkeley

Berkeley, CA

Service

2026	Program Committee , OOPSLA 2026: 47th ACM SIGPLAN Conference on Programming Language Design and Implementation	Oakland, CA
2025	Organizer , JuliaCon Global 2025: Sparse and Graph Computing in Julia Minisymposium	Pittsburgh, PA
2025	Program Committee , PLDI 2025: 46th ACM SIGPLAN Conference on Programming Language Design and Implementation	Seoul, South Korea
2025	Program Committee , CC25: ACM SIGPLAN 2025 International Conference on Compiler Construction (CC'25) co-located with CGO, PPoPP and HPCA	Las Vegas, NV
2022-2024	Organizer , Sparse Roofline Benchmark Working Group, a multi-university collaboration to standardize benchmarking of sparse linear algebra	Virtual
2022-2024	Editor , GraphBLAS BinSparse Binary Sparse File Format Standards Committee	Virtual
2023	Program Committee , DRAGSTERS (Distributions, Relational Algebra, Graphs, Semi-Rings, Tensors, and All That) workshop at PLDI conference	Orlando, FL
2022	Reviewer , Parallel Computing	N/A
2020	Reviewer , IEEE Transactions on Computers	N/A
2020	Reviewer , IEEE Transactions on Computers	N/A
2019	Reviewer , IEEE Transactions on Computers	N/A
2021	Reviewer , IEEE Transactions on Parallel and Distributed Systems	N/A
2020	Reviewer , IEEE Transactions on Parallel and Distributed Systems	N/A
2019	Reviewer , IEEE Transactions on Parallel and Distributed Systems	N/A

Experience

2022 NSF I-Corps Fall Cohort #2 - South Regional Node Program

Virtual

ENTREPRENEURIAL LEAD

Aug. 2022 - Nov. 2022

- Interviewed 100 potential customers to validate the market for tensor algebra compilers.

Sandia National Laboratory	Albuquerque, NM
CSGF PRACTICUM INTERN, SUPERVISOR: ERIK BOMAN	May 2019 – Aug 2019
• Discovered algorithms to reorganize sparse matrix nonzeros into dense blocks. Proposed the 1D-VBR sparse matrix format. Julia.	
MIT Julia Lab	Cambridge, MA
RESEARCH ASSISTANT	Sep 2016 – May 2019
• Worked with Professor Alan Edelman to develop Julia abstractions for scientific computing.	
• Developed prototype Julia implementation for the CLIMA earth and atmosphere modeling project.	
• Developed array programming infrastructure based on Julia base array interface.	
Los Alamos National Laboratory	Los Alamos, NM
RESEARCH INTERN, SUPERVISOR: HAI AH NAM	May 2016 – Aug 2016
• Parallelized a coupled cluster doubles nuclear quantum physics simulation to run on Wolf cluster. C++/MPI.	
NVIDIA	Santa Clara, CA
SOFTWARE ENGINEERING INTERN	June 2014 – Aug 2014
• Worked in a team to create a CPU profiler, intercepting dll calls and sampling using signal handlers. Created a small real-time system to handle stack traces and process them into various types of call graphs.	
• Fixed bugs, conducted testing, and wrote a test for cuda-gdb.	

Presentations

2025	"PyData/Sparse and Finch: extending sparse computing in the Python ecosystem" , Pydata Global 2025	Online
2025	"Finch: Sparse and Structured Tensor Programming with Control Flow" , Dependently Typed, PL Club at Georgia Tech	Atlanta, GA
2025	"Sparsity-Aware Autoscheduling for Numpy with Finch and Galley" , Sparse Workshop, PLDI	Seoul, South Korea
2025	"Finch.jl: Flexible and Efficient Sparse Tensor Programming!" , JuliaCon 2025	Pittsburgh, PA
2025	"Sparse and Structured Tensor Programming" , Oak Ridge National Laboratory, Computer Science and Mathematics Division	Oak Ridge, TN
2024	Panelist , Working Toward an Interface for Sparse BLAS, Supercomputing 2024	Atlanta, GA
2024	"binsparse: A Binary Sparse Tensor File Format" , HPC Graph Toolkits and the GraphBLAS Forum, Supercomputing 2024	Atlanta, GA
2024	"Sparse and Structured Array Programming" , BU Principles of Programming and Verification (POPV), BU	Boston, MA
2024	"Sparse and Structured Array Programming" , Google	Virtual
2024	"Finch: Sparse and Structured Array Programming" , Ph.D. Defense, MIT	Cambridge, MA
2024	"Finch: Sparse and Structured Array Programming" , Sammaniversary, MIT	Cambridge, MA
2024	"Compiling Control Flow in Sparse and Structured Array Programs" , Sparse Workshop, PLDI	Copenhagen, Denmark
2024	"Compilers for Sparse and Structured Array Programming" , International Workshop on Efficient Generative AI, University of Edinburgh	Edinburgh, UK
2024	"Finch: A Sparse and Structured Array Compiler in Julia" , Cambridge Area Julia Users Network	Cambridge, MA
2024	"Supporting the Sparse BLAS: File Formats and Benchmarks" , Sparse BLAS Workshop 2, University of Tennessee	Knoxville, TN
2023	"Sparse Compilers, Sparse Benchmarks" , Sparse BLAS Workshop, University of Tennessee	Knoxville, TN
2023	"Exploring the Design Space of Sparsity Through Compilers" , The Sparse Rooflines Group	Virtual
2023	"Exploring the Design Space of Sparsity Through Compilers" , RelationalAI Virtual Talk	Virtual
2023	"Finch: A Compiler for Sparse and Structured Data" , Stanford University	Stanford, CA
2023	"Finch: A Compiler for Sparse and Structured Data" , Lawrence Berkeley National Lab	Berkeley, CA
2023	"Finch: A Compiler for Sparse and Structured Data" , University of Washington PLSE Group	Seattle, WA
2022	"An Asymptotic Cost Model for Autoscheduling Sparse Tensor Programs" , ADA Symposium	Ann Arbor, Michigan
2021	"Contiguous Partitioning: Registers, Caches, and Distributed Memories" , DOE CSGF Review	Washington, D.C.
2021	"On Optimal Partitioning for Variable Block Row Format" , MIT CRIBB Seminar	Cambridge, MA
2018	"The Tensor Algebra Compiler (taco)" , CSAIL Alliances Annual Meeting	Cambridge, MA
2018	"For-Loops 2.0: Index Notation And The Future Of Array Compilers" , JuliaCon 2018	London, UK

Complete Publications

Insum: Sparse GPU Kernels Simplified and Optimized with Indirect Einsums.

ASPLOS 2026

J. WON, W. AHRENS, J. S. EMER, AND S. AMARASINGHE,

Mar. 2026

- To be published in ASPLOS '26: Proceedings of the 31st ACM International Conference on Architectural Support for Programming Languages and Operating Systems.

Binsparse: A Specification for Cross-Platform Storage of Sparse Matrices and Tensors

arXiv:2506.19175 [cs]

B. BROCK, W. AHRENS, H. ABBASI, T. DAVIS, J. KIM, J. KITCHEN, S. PATTY, I. VIRSHUP, E. WELCH

Jun. 2025

- Published in arXiv:2506.19175 [cs].
- Ongoing Specification: <https://github.com/GraphBLAS/binsparse-specification>

The Continuous Tensor Abstraction: Where Indices Are Real

OOPSLA 2025

J. WON, W. AHRENS, T. F. COLLIN, J. S. EMER, AND S. AMARASINGHE.

Apr. 2025

- Published in Proceedings of the ACM on Programming Languages, Volume 9, Issue OOPSLA2.

Finch: Sparse and Structured Tensor Programming with Control Flow

OOPSLA 2025

W. AHRENS, T. F. COLLIN, R. PATEL, K. DEEDS, C. HONG, AND S. AMARASINGHE.

Apr. 2025

- Published in Proceedings of the ACM on Programming Languages, Volume 9, Issue OOPSLA1.

SySTeC: A Symmetric Sparse Tensor Compiler

CGO 2025

R. PATEL, W. AHRENS, AND S. AMARASINGHE.

Mar. 2025

- Published in Proceedings of the 23rd ACM/IEEE International Symposium on Code Generation and Optimization.

Interface for Sparse Linear Algebra Operations.

arXiv:2411.13259 [cs]

A. ABDELFATTAH, W. AHRENS, H. ANZT, ET AL.

Nov. 2024

- Published in arXiv:2411.13259 [cs].

Galley: Modern Query Optimization for Sparse Tensor Programs.

SIGMOD 2025

K. DEEDS, W. AHRENS, M. BALAZINSKA, AND D. SUCIU.

Jun. 2025

- Published in Proceedings of the ACM on Management of Data, Volume 3, Issue 3.

SySTeC: A Symmetric Sparse Tensor Compiler

CGO 2025

R. PATEL, W. AHRENS, AND S. AMARASINGHE.

Mar. 2025

- Published in Proceedings of the 23rd ACM/IEEE International Symposium on Code Generation and Optimization, in CGO 25.

Mechanised Hypersafety Proofs about Structured Data

PLDI 2024

V. GLADSHTEIN, Q. ZHAO, W. AHRENS, S. AMARASINGHE, AND I. SERGEY.

Jun. 2024

- Published in Proceedings of the ACM on Programming Languages, Volume 8, Issue PLDI, p. 173:647-173:670. Jun. 2024

Looplets: A Language for Structured Coiteration.

CGO 2023

W. AHRENS, D. DONENFELD, F. KJOLSTAD, AND S. AMARASINGHE.

Feb. 2023

- Published in Proceedings of the 21st ACM/IEEE International Symposium on Code Generation and Optimization, in CGO 2023.

Autoscheduling For Sparse Tensor Algebra With An Asymptotic Cost Model.

PLDI 2022

W. AHRENS, F. KJOLSTAD, AND S. AMARASINGHE.

Jun. 2022

- Published in Proceedings of the 43rd ACM SIGPLAN International Conference on Programming Language Design and Implementation.

Contiguous Graph Partitioning For Optimal Total Or Bottleneck Communication.

arXiv:2007.16192 [cs]

W. AHRENS.

Jun. 2021

- Published in arXiv:2007.16192 [cs].

Algorithms for Efficient Reproducible Floating Point Summation.

ACM Trans. Math. Softw.

W. AHRENS, J. DEMMEL, AND H. D. NGUYEN.

Jul. 2020

- Published in ACM Transactions on Mathematical Software, vol. 46, no. 3, p. 22:1-22:49, Jul. 2020.

Brief Announcement: Sparse Tensor Transpositions.

SPAA 2020

S. MUELLER, W. AHRENS, S. CHOU, F. KJOLSTAD, AND S. AMARASINGHE.

2020

- Published in Proceedings of the 32nd ACM Symposium on Parallelism in Algorithms and Architectures (SPAA), 2020, pp. 559-561.

On Optimal Partitioning For Sparse Matrices In Variable Block Row Format.

arXiv:2005.12414 [cs]

W. AHRENS AND E. G. BOMAN.

May 2020

- Published in arXiv:2005.12414 [cs].

A Parallel Fill Estimation Algorithm for Sparse Matrices and Tensors in Blocked Formats.

MIT Thesis

W. AHRENS.

2019

- Thesis submitted to the Massachusetts Institute of Technology in 2019.

Tensor Algebra Compilation with Workspaces.

CGO 2019

F. KJOLSTAD, W. AHRENS, S. KAMIL, AND S. AMARASINGHE.

2019

- Published in 2019 IEEE/ACM International Symposium on Code Generation and Optimization (CGO), 2019, pp. 180-192.

LATE Ain't Earley: A Faster Parallel Earley Parser.

arXiv:1807.05642 [cs]

W. AHRENS, J. FESER, AND R. HUI.

Jul. 2018

- Published in arXiv:1807.05642 [cs].

A Fill Estimation Algorithm for Sparse Matrices and Tensors in Blocked Formats.

IPDPS 2018

W. AHRENS, H. XU, AND N. SCHIEFER.

2018

- Published in 2018 IEEE International Parallel and Distributed Processing Symposium (IPDPS), 2018, pp. 546-556.

Parallel Compact Hash Algorithms for Computational Meshes.

SIAM J. Sci. Comput.

R. TUMBLIN, W. AHRENS, S. HARTSE, AND R. ROBEY.

Jan. 2015

- Published in SIAM Journal on Scientific Computing, vol. 37, no. 1, pp. C31-C53, Jan. 2015.